

AZURE ML: FLIGHT DELAY MODEL

Jess Chang 張傳忠

jechang@microsoft.com

Architect

Microsoft Technology Center

2014.12

Overview

- Data Insights
 - ▶ Inline visualizer
 - ▶ Descriptive statistics
- Data Scrubbing
 - ▶ Set missing values
 - ▶ Scrub missing values
 - ▶ Project subset of columns
 - ▶ Edit metadata
- Data Transformation (Part one)
 - ▶ Apply math operations
 - ▶ Remove duplicate rows
 - ▶ Join two datasets
 - ▶ Split data
- Data Transformation (Part two)
 - ▶ Transform data using R script

Sample Application - Predicting Flight Delay

- Problem
 - ▶ Approx. 20% of flights are delayed or cancelled every year.
 - ▶ Many factors affect delays: weather, mechanical issues, air traffic control, etc.
- Goal
 - ▶ Leverage past flight and weather data to predict future flight delays.
 - ▶ Reference: [*Predicting Flight Delays*](#), Dieterich Lawson and William Castillo, Stanford University
- Data sources
 1. Airline On-time Performance dataset
 - ▶ Source: [Bureau of Transportation Statistics \(BTS\)](#)
 2. Weather Observations (ISD-Lite) dataset
 - ▶ Source: [National Oceanic and Atmospheric Administration \(NOAA\)](#)
 - ▶ Dataset: Quality Controlled Local Climatological Data ([FTP](#))

A sample from July to October 2013 will be used in this exercise.

Flight On-Time Performance Dataset

| Year | Month | DayofMonth | DayOfWeek | Carrier | OriginAirportID | DestAirportID | CRSDepTime | DepDelay | DepDel15 | CRSArrTime | ArrDelay | ArrDel15 | Cancelled |
|------|-------|------------|-----------|---------|-----------------|---------------|------------|----------|----------|------------|----------|----------|-----------|
| 2013 | 7 | 16 | 2 | DL | 13487 | 14747 | 2155 | -1 | 0 | 2336 | 5 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 12889 | 13487 | 1555 | -6 | 0 | 2057 | -7 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 11278 | 10397 | 1600 | -5 | 0 | 1752 | -19 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 13851 | 10397 | 600 | -3 | | 904 | 4 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 14747 | 12 | 2330 | 49 | | 736 | 40 | 1 | 0 |
| 2013 | 7 | 16 | 2 | DL | | | 1735 | -4 | | 2108 | -41 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | | | 1656 | 33 | | 1831 | 17 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | | | 659 | -7 | | 837 | -28 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | | | 805 | -2 | | 859 | -25 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | | | 1005 | -6 | | 1650 | -8 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 12953 | 10397 | 700 | 112 | | 930 | 108 | | 0 |
| 2013 | 7 | 16 | 2 | DL | 11433 | 12953 | 1725 | | | 14 | | | |
| 2013 | 7 | 16 | 2 | DL | 13495 | 12892 | 720 | | | 17 | | | |
| 2013 | 7 | 16 | 2 | DL | 12889 | 13487 | 720 | | | 22 | | | |
| 2013 | 7 | 16 | 2 | DL | 13487 | 12889 | 1750 | | | 11 | -2 | | |
| 2013 | 7 | 16 | 2 | DL | 12889 | 12892 | 620 | | | 30 | | | |
| 2013 | 7 | 16 | 2 | DL | 12889 | 10397 | 715 | -7 | 0 | 1415 | -1 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 10397 | 12892 | 940 | -2 | 0 | 1110 | -11 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 15304 | 10397 | 1445 | -5 | 0 | 1615 | -5 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 14869 | 14747 | 2155 | -5 | 0 | 2300 | 31 | 1 | 0 |
| 2013 | 7 | 16 | 2 | DL | 15304 | 12892 | 1930 | -8 | 0 | 2125 | -17 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 13487 | 12892 | 1135 | 2 | 0 | 1321 | -17 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 12892 | 12173 | 1442 | 4 | 0 | 1723 | 2 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 11433 | 10529 | 720 | -4 | 0 | 900 | -9 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 10529 | 11433 | 941 | -4 | 0 | 1129 | -12 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 10397 | 14100 | 1117 | -4 | 0 | 1320 | -5 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 14100 | 10397 | 1415 | -5 | 0 | 1616 | -12 | 0 | 0 |
| 2013 | 7 | 16 | 2 | DL | 14869 | 13487 | 2016 | 6 | 0 | 2346 | 0 | 0 | 0 |

Carrier is a categorical field

Hour and minutes concatenated in one field

OriginAirportID and DestAirportID are categorical fields with numeric values

Multiple target leaks: DepDelay, DelDel15, ArrDelay, Cancelled

Assume ArrDel15 is the target to be predicted



Windows Azure
Machine Learning

Sample flight data

NOAA Weather Observations Dataset

| Year | Month | Day | AirportID | Time | TimeZone | SkyCondition | Visibility | WeatherType | DryBulbFarenheit | DryBulbCelsius | WetBulbFarenheit | WetBulbCelsius | DewPointFarenheit | DewPointCelsius | RelativeHumidity |
|-----------------------------------|-------|-----|-----------|------|----------|-----------------|------------|-------------|------------------|----------------|------------------|----------------|-------------------|-----------------|------------------|
| 2013 | 7 | 1 | 14843 | 56 | -4 | FEW020 SCT035 | 10 | | 78 | 25.6 | 75 | 23.6 | 73 | 22.8 | 85 |
| 2013 | 7 | 1 | 14843 | 156 | -4 | FEW035 | 10 | | 78 | 25.6 | 74 | 23.2 | 72 | 22.2 | 82 |
| 2013 | 7 | 1 | 14843 | 256 | -4 | FEW050 | 10 | | 78 | 25.6 | 75 | 23.6 | 73 | 22.8 | 85 |
| 2013 | 7 | 1 | 14843 | 356 | -4 | FEW055 SCT070 | 10 | | 78 | 25.6 | 75 | 23.6 | 73 | 22.8 | 85 |
| 2013 | 7 | 1 | 14843 | 456 | -4 | FEW050 | 10 | | 77 | 25 | 74 | 23.4 | 73 | 22.8 | 88 |
| 2013 | 7 | 1 | 14843 | 556 | -4 | FEW055 SCT065 | 10 | | 78 | 25.6 | 75 | 23.6 | 73 | 22.8 | 85 |
| Time is in UTC, not local time | | | | | | | | | | | | | | | |
| 2013 | 7 | 1 | 14843 | 656 | -4 | N020 SCT042 SCT | 10 | | 78 | 25.6 | 75 | 24 | 74 | 23.3 | 88 |
| 2013 | 7 | 1 | 14843 | 756 | -4 | N034 SCT032 SCT | 10 | | 81 | 27.2 | 77 | 24.8 | 75 | 23.9 | 82 |
| 2013 | 7 | 1 | 14843 | 856 | -4 | N034 SCT032 SCT | 9 | | 85 | 29.4 | 77 | 25.1 | 74 | 23.3 | 70 |
| 2013 | 7 | 1 | 14843 | 956 | -4 | FEW022 S | 9 | | 86 | 30 | 78 | 25.6 | 75 | 23.9 | 70 |
| 2013 | 7 | 1 | 14843 | 1056 | -4 | FEW025 SCT | 9 | | 87 | 30.6 | 78 | 25.8 | 75 | 23.9 | 68 |
| 2013 | 7 | 1 | 14843 | 1156 | -4 | 025 SCT032CB | | TS | 84 | 29 | 77 | 25 | 74 | 23.3 | 72 |
| 2013 | 7 | 1 | 14843 | 1215 | -4 | T025 S | | | 84 | 29 | 76 | 24.6 | 73 | 23 | 70 |
| 2013 | 7 | 1 | 14843 | 1256 | -4 | V025 S | | | 81 | 27.2 | 77 | 24.8 | 75 | 23.9 | 82 |
| 2013 | 7 | 1 | 14843 | 1356 | -4 | T028 S | | | 83 | 28.3 | | 24.4 | 73 | 22.8 | 72 |
| 2013 | 7 | 1 | 14843 | 1456 | -4 | V038 S | | | 84 | 28.9 | | | 74 | 23.3 | 72 |
| 2013 | 7 | 1 | 14843 | 1556 | -4 | V036 SCT032CB | | | 83 | 28.3 | | | 73 | 22.8 | 72 |
| 2013 | 7 | 1 | 14843 | 1656 | -4 | FEW042 BKN100 | 9 | | 83 | 28.3 | | | 74 | 23.3 | 74 |
| 2013 | 7 | 1 | 14843 | 1756 | -4 | FEW055 BKN100 | 9 | | 83 | 28.3 | | | 74 | 23.3 | 74 |
| 2013 | 7 | 1 | 14843 | 1856 | -4 | V038 SCT075 BKN | 9 | | 82 | 27.8 | 76 | 24.3 | 73 | 22.8 | 74 |
| 2013 | 7 | 1 | 14843 | 1956 | -4 | FEW055 | 10 | | 81 | 27.2 | 75 | 24.1 | 73 | 22.8 | 77 |
| 2013 | 7 | 1 | 14843 | 2056 | -4 | FEW045 SCT065 | 10 | | 81 | 27.2 | 75 | 24.1 | 73 | 22.8 | 77 |
| 2013 | 7 | 1 | 14843 | 2156 | -4 | FEW025 SCT055 | 10 | | 80 | 26.7 | 75 | 23.9 | 73 | 22.8 | 79 |
| 2013 | 7 | 1 | 14843 | 2256 | -4 | FEW025 SCT055 | 10 | | 80 | 26.7 | 76 | 24.3 | 74 | 23.3 | 82 |
| 2013 | 7 | 1 | 14843 | 2356 | -4 | FEW025 SCT060 | 10 | | 79 | 26.1 | 76 | 24.1 | 74 | 23.3 | 85 |
| 2013 | 7 | 2 | 14843 | 56 | -4 | FEW022 BKN070 | 10 | | 79 | 26.1 | 76 | 24.1 | 74 | 23.3 | 85 |
| 2013 | 7 | 2 | 14843 | 156 | -4 | FEW040 | 10 | =RA | 79 | 26.1 | 75 | 23.8 | 73 | 22.8 | 82 |
| 2013 | 7 | 2 | 14843 | 256 | -4 | FEW040 | 10 | =RA | 78 | 25.6 | 75 | 23.6 | 73 | 22.8 | 85 |
| 2013 | 7 | 2 | 14843 | 356 | -4 | FEW022 SCT040 | 10 | | 78 | 25.6 | 75 | 24 | 74 | 23.3 | 88 |

Windows Azure

Machine Learning

Sample weather data

Prepare Data Transformation Experiment

- Download the two sample datasets for July-October 2013 locally
 - ▶ [Flight On-Time Performance](#)
 - ▶ [NOAA Weather Observations](#)
- Upload the datasets to Azure ML
 1. Go to <https://studio.azureml.net>
 2. Click on **+New** to create a new Dataset
 3. Click on **Dataset → From Local File**
 4. Upload the Flight On-Time Performance local file.
 5. Repeat to upload the NOAA Weather Observations local file.
- Start new experiment
 1. Go to <https://studio.azureml.net> → Click on **Experiments**
 2. Click on **+New →Experiment** to create a new experiment
 3. Double click on “Untitled” on top of the screen and type in “**Flight delay prediction**”
 4. Expand **Saved Datasets** in the left panel.
 5. Drag the dataset **Flight_On_Time_Performance_July_October_2013.csv** to the experiment.
 6. Drag the dataset **NOAA_Weather_July_October_2013.csv** to the experiment.

Process the Flight On time Performance Dataset

1. Ingest and set missing values to ‘?’

- ▶ Module: Convert to Dataset
- ▶ SetMissingValues → ?

2. Explore the data

- ▶ Click output node → Visualize
- ▶ Module: Descriptive Statistics

3. Remove target leak fields from dataset

- ▶ Module: Project Columns
- ▶ Begin With All Columns
- ▶ Exclude columns [DepDelay, DepDel15, ArrDelay, Cancelled]

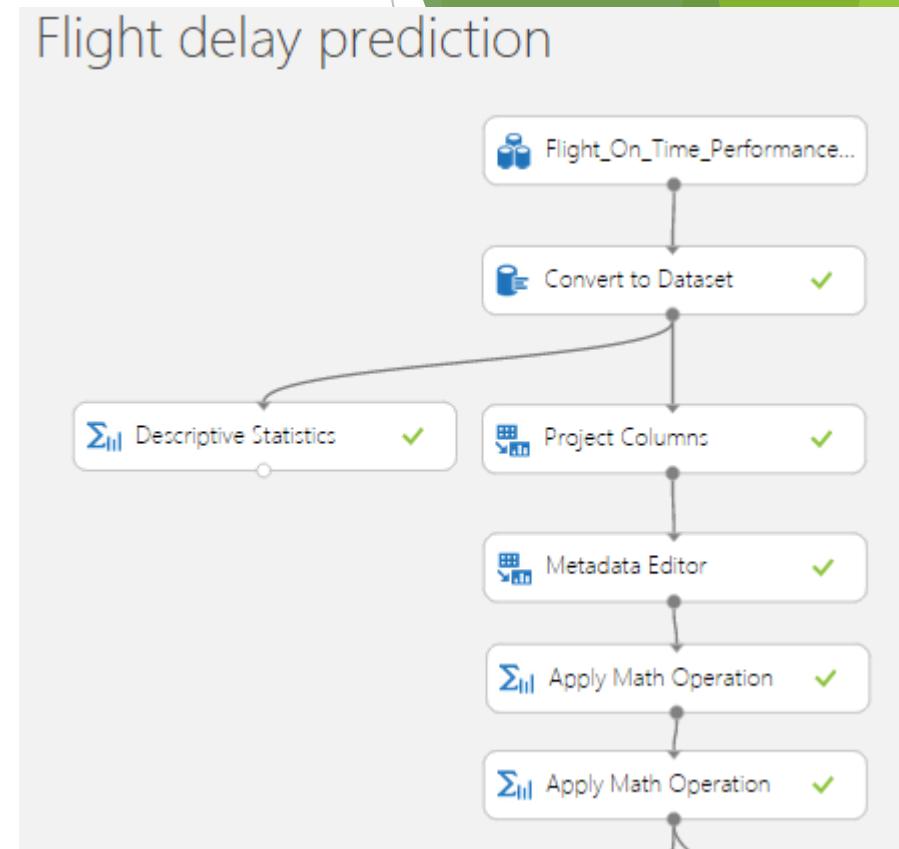
4. Edit metadata to indicate **categorical** fields

- ▶ Module: Metadata Editor
- ▶ Begin With NO Columns
- ▶ Include columns [Carrier, OriginAirportID, DestAirportID] to **Categorical**

5. Extract *hour* from time fields [CRSDepTime and CRSArrTime]

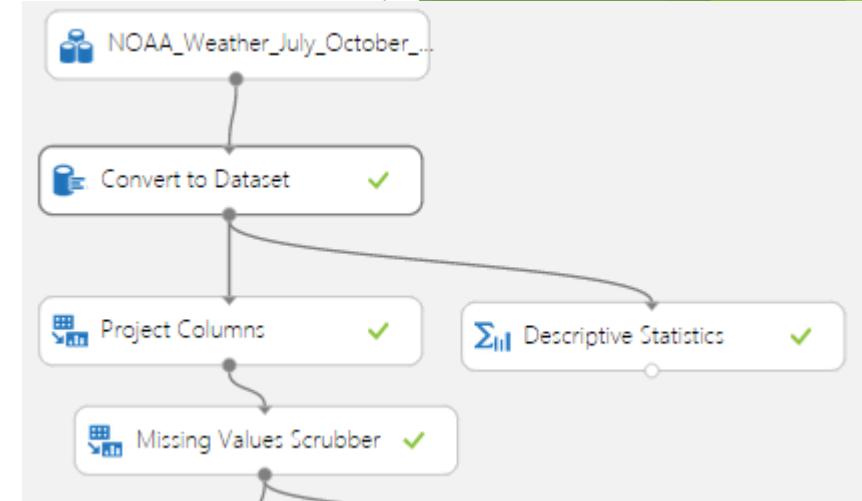
- ▶ Module: Apply Math Operation
- ▶ Step 1: Operations → Divide *time* field by 100 → Begin With NO Columns → Include columns [CRSDepTime and CRSArrTime]
- ▶ Module: Apply Math Operation
- ▶ Step 2: Rounding → Floor operation → Begin With NO Columns → Include columns [CRSDepTime and CRSArrTime]
- ▶ Output mode: **Inplace**

Flight delay prediction



Process the NOAA Weather Observations Dataset

1. Ingest and set missing values to 'M'
 - ▶ Module: Convert to Dataset
 - ▶ SetMissingValues → M
2. Explore the data
 - ▶ Click output node → Visualize
 - ▶ Module: Descriptive Statistics
3. Remove unnecessary fields from dataset
 - ▶ Module: Project Columns
 - ▶ Begin With All Columns
 - ▶ Exclude columns names [WetBulbFarenheit, WetBulbCelsius, ValueForWindCharacter, StationPressure, PressureTendency, PressureChange, SeaLevelPressure]
 - ▶ Exclude columns types [String]
4. Remove rows with missing values
 - ▶ Module: Missing Values Scrubber
 - ▶ For missing values → Remove entire row
 - ▶ Cols with all MV → KeepColumns
 - ▶ MV indicator column → DoNotGenerate



Process the NOAA Weather Observations Dataset (ctd.)

1. Extract *hour* from time field [Time]

- ▶ Module: Apply Math Operation
- ▶ Step 1: Operations → Divide *time* field by 100 → **Begin With NO Columns** → **Include** columns [Time]
- ▶ Module: Apply Math Operation
- ▶ Step 2: Rounding → Ceiling operation → **Begin With NO Columns** → **Include** columns [Time]
- ▶ Output mode: **Inplace**

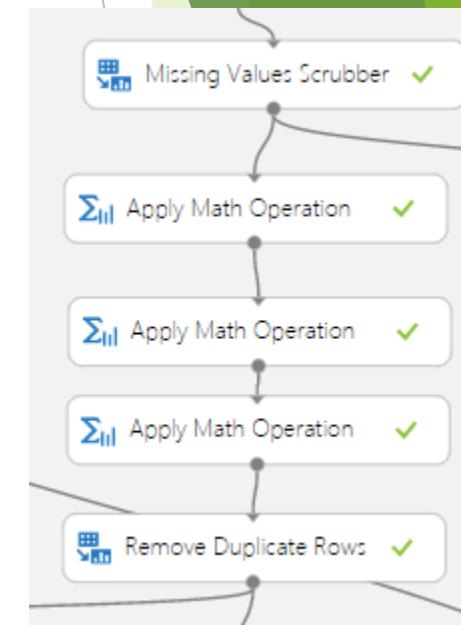
2. Adjust time from UTC to local time using TimeZone field

- ▶ Module: Apply Math Operation
- ▶ Operations → Subtract → ColumnSet
- ▶ Operation argument → **Begin With NO Columns** → Select column [TimeZone]
- ▶ Column set → **Begin With NO Columns** → Select column [Time]
- ▶ Output mode: **Append**

❖ Save Experiment and Run

1. Remove duplicate rows to keep one weather observation per hour

- ▶ Module: Remove Duplicate Rows
- ▶ Key column selection filter expression → **Begin With NO Columns** → Select columns [Year, Month, Day, AirportID, Subtract(Time_TimeZone)]



Join Flight and Weather Datasets + Split

1. Join the two datasets to add weather information at departure airport

**Begin With NO Columns*

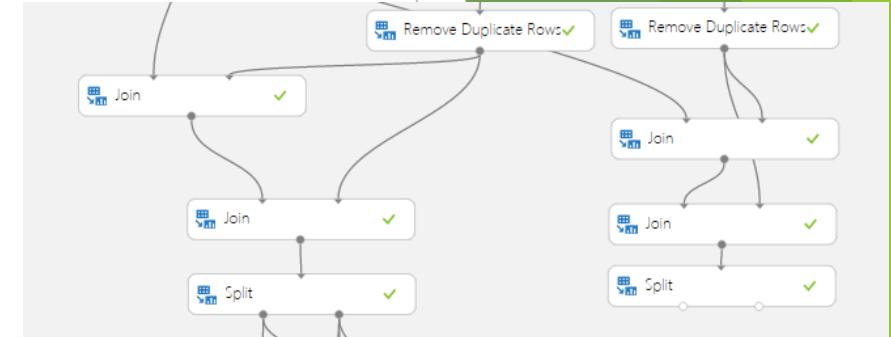
- ▶ Module: Join
- ▶ Join key columns for L → Select columns [Year, Month, DayOfMonth, **OriginAirportID**, CRSDepTime]
- ▶ Join key columns for R → Select columns [Year, Month, Day, AirportID, Subtract(Time_TimeZone)]
- ▶ Join type → Inner Join
- ▶ Uncheck “Keep right key columns in joined table”

1. Join the two datasets to add weather information at arrival airport

- ▶ Repeat (1) with the following left key columns
- ▶ Join key columns for L → Select columns [Year, Month, DayOfMonth, **DestAirportID**, CRSDepTime]

2. Split the final data to separate October 2013 data

- ▶ Module: Split
- ▶ Splitting mode → Relative Expression
- ▶ Relational expression → Enter the value \\"Month" < 10



Transform Data Using an R Script

- Motivation
 - ▶ Perform multiple transformations in one step
 - ▶ Complement Azure ML Studio with additional operations, functions, packages, etc.
- Exercise: Replace NOAA Weather Observations data processing with an R script
 - ▶ Extract *hour* from the Time field
 - ▶ Adjust weather time from UTC to local time using DateTime operations
 - ▶ Append the adjusted Month/Day/Hour fields to input dataset
 - ▶ Project required output columns and re-order them
 - ▶ R script code in next slide
- Run the R script code in Azure ML Studio
 - ▶ Module: **Execute R Script**
 - ▶ Copy the R script from your favorite editor (e.g., Rstudio)
 - ▶ Paste the code in the designated *Script* area in the module Properties panel
 - ▶ Random Seed: 42
- ❖ Save Experiment and Run

Weather Data Transformation R Script

```
# Map input port to variable  
  
dataset <- maml.mapInputPort(1) # class: data.frame  
  
  
Year = dataset$Year  
Month = dataset$Month  
Day = dataset$Day  
Time = dataset$Time  
Hour = ceiling(Time / 100)  
Timezone = dataset$TimeZone  
  
  
# Number of rows to process  
n = nrow(dataset)  
  
  
# Concatenate date fields and apply time zone difference  
  
fulldate = lapply(1:n, function(i) as.POSIXlt(sprintf("%4d-%02d-%02d %02d:00:00", Year[i], Month[i], Day[i], Hour[i]), tz = "UTC") - Timezone[i] * 3600)  
adjustdate = do.call(c,fulldate)  
  
  
# Extract the adjusted month, day, and hour - Should adjust year too for general case  
AdjustedMonth = as.POSIXlt(adjustdate)$mon + 1  
AdjustedDay = as.POSIXlt(adjustdate)$mday  
AdjustedHour = as.POSIXlt(adjustdate)$hour  
  
  
# Extract other columns  
AirportID = dataset$AirportID  
Weather = dataset[,7:14]  
  
  
# Construct output data frame and send to the output Dataset port  
data.set = cbind(Year, AdjustedMonth, AdjustedDay, AirportID, AdjustedHour, Weather)  
maml.mapOutputPort("data.set");
```

Join Flight and Weather (R executed) Datasets + Split

- Remove duplicate rows to keep one weather observation per hour

*Begin With NO Columns

- ▶ Module: Remove Duplicate Rows
- ▶ Key column selection filter expression → Select columns [Year, AdjustedMonth, AdjustedDay, AirportID, AdjustedHour]

1. Join the two datasets to add weather information at departure airport

- ▶ Module: Join
- ▶ Join key columns for L → Select columns [Year, Month, DayOfMonth, OriginAirportID, CRSDepTime]
- ▶ Join key columns for R → Select columns [Year, AdjustedMonth, AdjustedDay, AirportID, AdjustedHour]
- ▶ Join type → Inner Join
- ▶ Uncheck “Keep right key columns in joined table”

2. Join the two datasets to add weather information at arrival airport

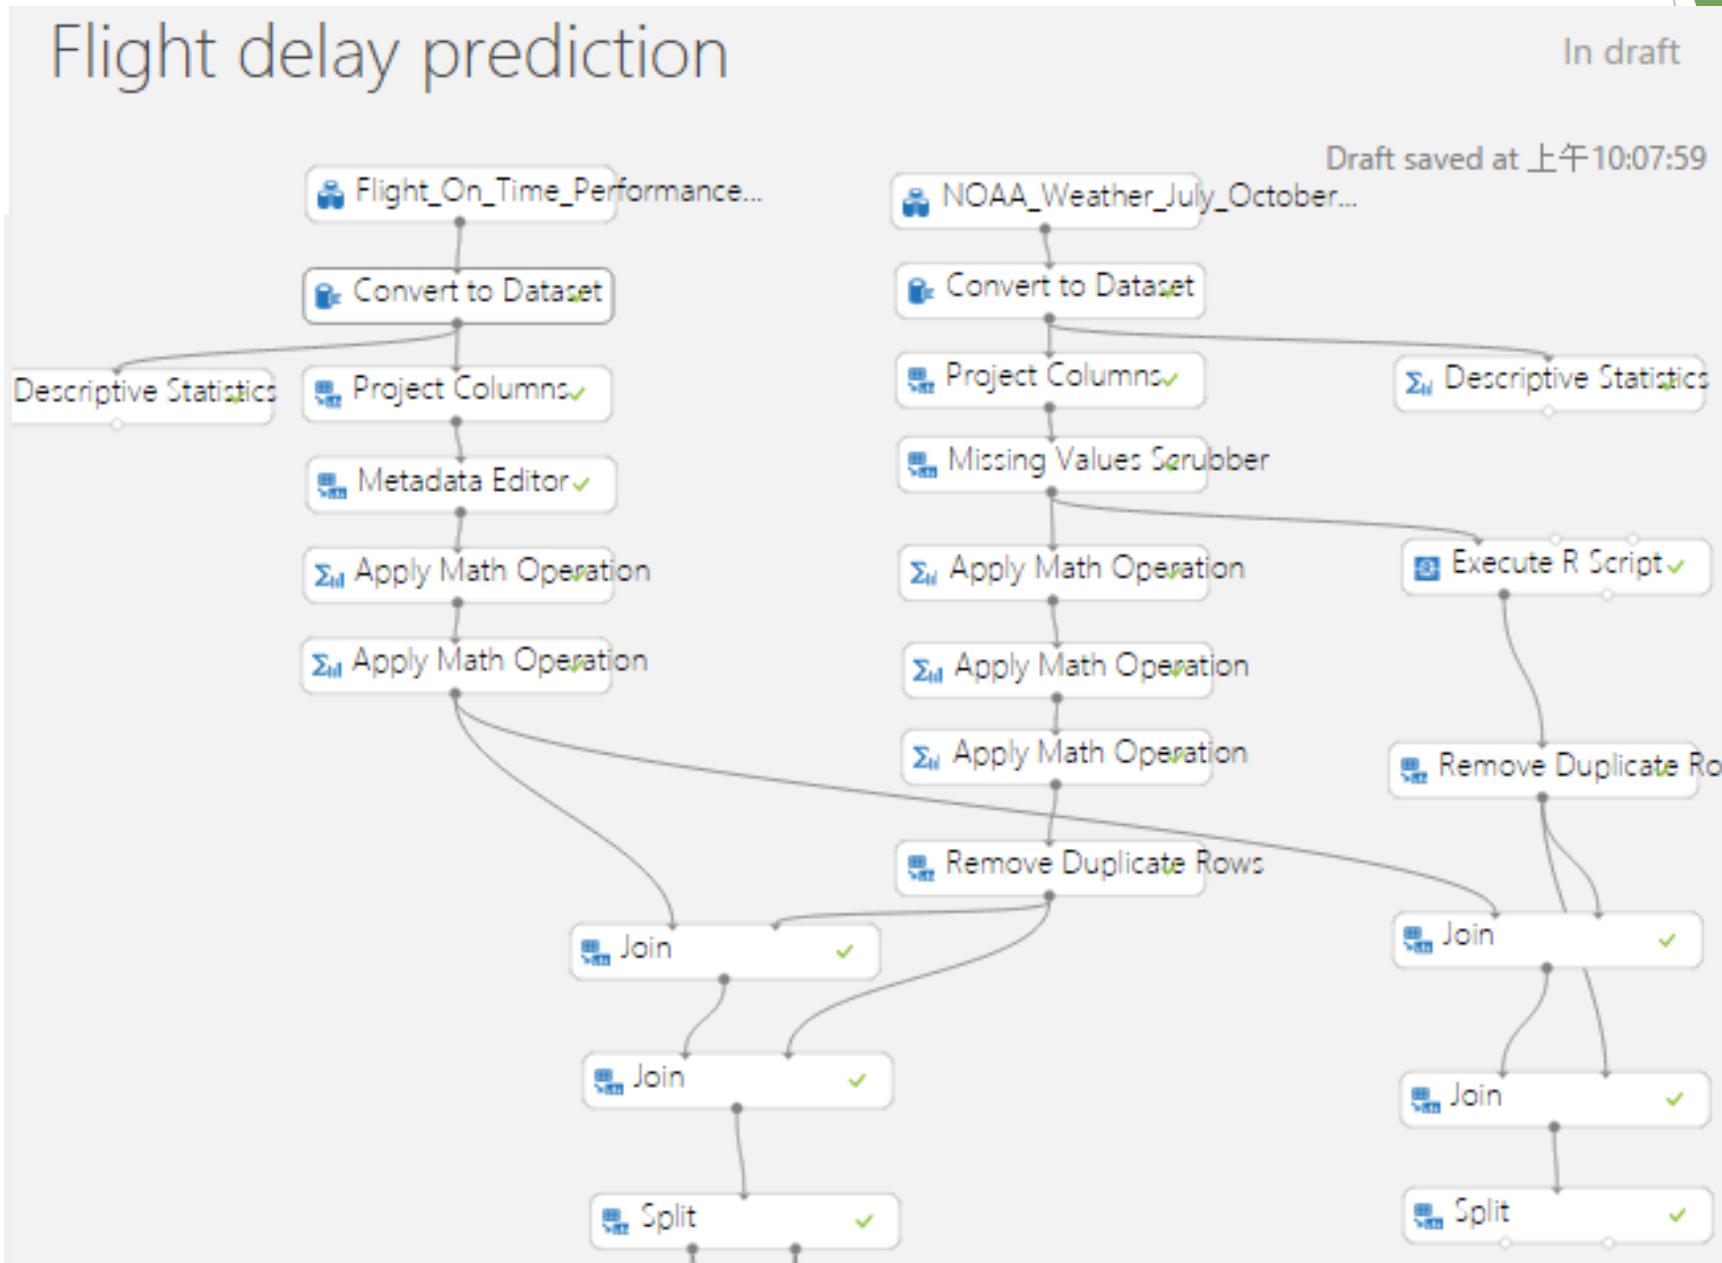
- ▶ Repeat (1) with the following left key columns
- ▶ Join key columns for L → Select columns [Year, Month, DayOfMonth, DestAirportID, CRSDepTime]

3. Split the final data to separate October 2013 data

- ▶ Module: Split
- ▶ Splitting mode → Relative Expression
- ▶ Relational expression → Enter the value \\"Month" < 10

Data Transformation Experiment

Flight delay prediction



Two-Class Boosted Decision Tree

1. Adopt Two-Class Boosted Decision Tree Module

- ▶ Maximum number of leaves per tree: **20**
- ▶ Minimum number of samples per leaf node: **10**
- ▶ Learning rate: **0.2**
- ▶ Number of trees constructed: **100**
- ▶ Allow unknown categorical levels

2. Optimize parameter settings: [Random sweep]

- ▶ Module: **Sweep Parameters**
- ▶ Maximum number of runs on random sweep: **10**
- ▶ Label column: Include Column names: **ArrDel15**
- ▶ Metric for measuring performance for classification: **Accuracy**
- ▶ Metric for measuring performance for regression: **Mean absolute error**
- ▶ Input Untrained Model(Left most): **Output of Two-Class Boosted Decision Tree**
- ▶ Input Training Dataset(Central): **Output of Split dataset1**
- ▶ Input Validation Dataset(Right most): **Output of Split dataset2**

3. Score a trained classification regression model

- ▶ Module: **Score Model**
- ▶ Input Trained Model(Left): **Output of Sweep parameters Trained best model(Right)**
- ▶ Input Dataset(Right): **Output of Split dataset2**

Two-Class Bayes Point Machine

1. Adopt Two-Class Bayes Point Machine Module

- ▶ Number of training iterations: 10
- ▶ Include bias
- ▶ Allow unknown values in categorical feature

2. Train a previously created classification or regression model

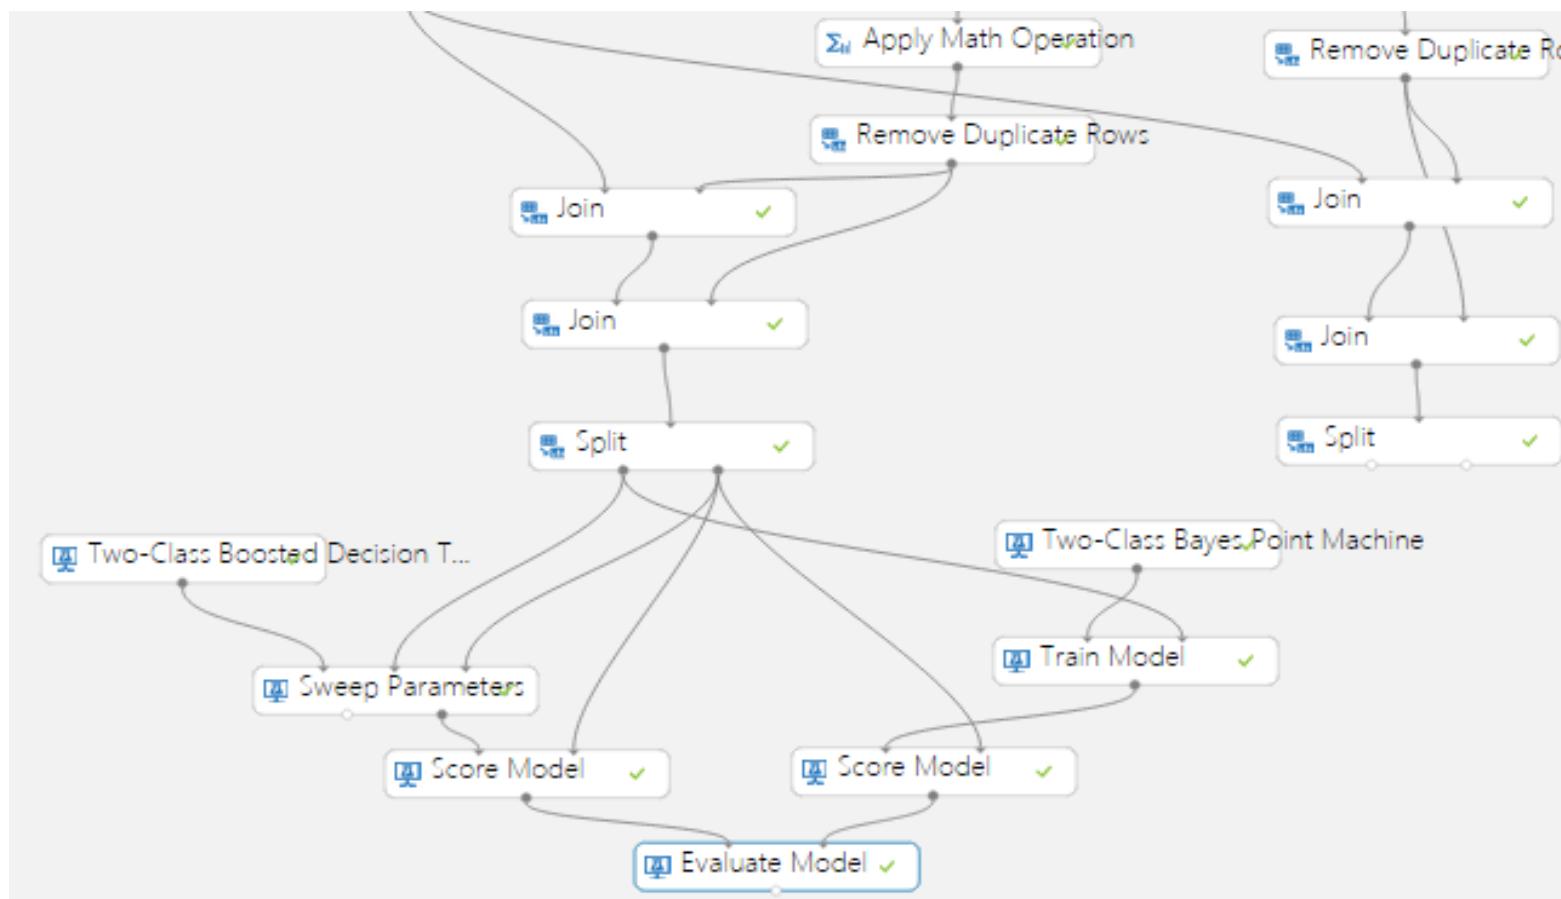
- ▶ Module: **Train Model**
- ▶ Include Label column: **[ArrDel15]**
- ▶ Input Untrained Model(Left): **Output of Two-Class Bayes Point Machine**
- ▶ Input Dataset(Right): **Output of Split dataset1**

3. Score a trained classification regression model

- ▶ Module: **Score Model**
- ▶ Input Trained Model(Left): **Output of Train Model**
- ▶ Input Dataset(Right): **Output of Split dataset2**

Evaluate a scored classification regression model

- ▶ Module: Evaluate Model
- ▶ Input scored dataset1: Score Model of Two-Class Boosted Decision Tree
- ▶ Input scored dataset2 to compare: Score Model of Two-Class Boosted Decision Tree



AZURE ML: FLIGHT DELAY MODEL